

# Predicting Market Volatility using Hybrid Models

25HR06

Abhinav Kumar	2301AI50
Swagatam Pati	2301AI28
Soumyakanta Panda	2301CS86
Kalp Alpesh Shah	2301CS93
Anurag Nath	2301CS07



# 🌀 volatility ?

understanding market volatility dynamics

- > *realized volatility* quantifies the degree of actual variability in asset prices over a specified period ... how much prices move up or down over time
- > big swings = high volatility  
small moves = low volatility ... tells you how unpredictable the market is
- > usually calculated from historical intraday or daily returns
- > typically computed using squared returns and often annualized
- > differs from *implied volatility*  
... derived from option prices and reflects market expectations



# applications ?

understanding applications of realized volatility

- > realized volatility provides a backward-looking, data-driven view of market uncertainty ... helps measure how unstable a stock or market has been
- > serves as a vital tool in quantitative finance for risk assessment, portfolio optimization, and pricing of derivatives
- > valuable in high-frequency trading and market micro-structure analysis, where precise measurement of short-term price dynamics is crucial
- > accurate estimation of this metric enhances model reliability and financial decision-making

## ⌚ problem formulation

- > **objective**  
predict short-term realized volatility using high-frequency limit order book [LOB] data
- > **challenges**
  - market micro-structure noise in high-frequency data
  - need for fine-grained feature engineering [e.g. WAP, order flow]
  - temporal dynamics and stock-specific behavior variation
- > **goal**  
design a model pipeline leveraging LightGBM, MLP, and CNN models to capture spatial and temporal dependencies
- > **output**  
a numerical prediction of volatility for each stock-time pair in the test set.

## ⌚ traditional models

- > **GARCH**  
good at modeling time-varying volatility  
but inadequate for high-frequency granularity
- > **ARIMA**  
effective for trend/seasonality but lacks  
predictive power for nonlinear volatility dynamics

## ⌚ modern approaches

- > **LSTM (long short-term memory)**  
captures sequential dependencies but suffers  
from long training times and over-fitting in  
small datasets
- > **CNNs for time series**  
shown to capture short-term patterns efficiently
- > **Gradient Boosting (LightGBM)**  
strong tabular learner for engineered features

# ⌚ dataset

## > source

*Optiver Realized Volatility Prediction* [@kaggle]

<https://www.kaggle.com/competitions/optiver-realized-volatility-prediction>

## > content

- book data (limit order book)
    - “bid/ask prices” and “sizes” for top 10 levels timestamped within each time bucket
  - trade data
    - - executed trades: “price”, “size”, “number\_of\_orders”
- 44 csv files ; 1 file for each stock  
1 file  
ask\_price, bid\_price, ask\_size, bid\_size, time\_id  
(1 cent)

## > target variable

- realized volatility per “stock\_id” and “time\_id”.
- calculated from log returns of WAP [weighted average price]

## > challenges

- high dimensionality due to LOB levels
- missing values and irregular time intervals
- stock-specific patterns vary drastically

# ⌚ preprocessing

- > data cleaning
  - removed rows with missing / invalid values in key columns [“prices”, “sizes”]
- > timestamp alignment
  - converted all relative timestamps to a consistent scale
  - grouped events by “time\_id” buckets [10-minute intervals]
- > price normalization
  - multiplied book price levels by estimated tick size to recover real-world prices

## ... tick size recovery

$$\text{tick}_i = \min(\text{diffs} \in \{p_j - p_{j-1} \mid p_j > p_{j-1} > 0\})$$

- > purpose
  - raw book prices are normalized [unit-less]
  - recovering tick size helps map them back to actual monetary values
- > method
  - for each stock:
    - - collected all unique non-zero level-1 prices [“ask\_price”, “bid\_price1”]
    - - sorted them and computed differences between adjacent prices
    - - used the minimum meaningful positive difference [above 1e-8] as the tick size



# time-order reconstruction

## .> objective

- reconstruct a global ordering of “time\_ids” across all stocks based on similarity
- enables time-aware models like CNNs and avoids temporal leakage in cross-validation

## 1> pivot table creation

- constructed a matrix of shape:  
 $(time\_id) \times (stock\_id)$  where each entry holds the end-of-bucket WAP
- missing entries filled using:
  - - forward fill (ffill), backward fill (bfill)

$$WAP = \frac{P_{ask} \cdot Q_{bid} + P_{bid} \cdot Q_{ask}}{Q_{ask} + Q_{bid}}$$

## 2> PCA (principal component analysis) for denoising

- applied PCA to reduce dimensionality
- captured dominant patterns in WAP movements

## 3> t-SNE for embedding

- applied t-distributed stochastic neighbor embedding (t-SNE) on PCA output:
  - - reduced dimensionality to 1D
  - - preserved temporal proximity of similar “time\_ids”
- - *perplexity=30, learning\_rate=200, n\_iter=1000*

## 4> time index mapping

- - sorted the embedded 1D values to produce a ranked ordering of all “time\_ids”



# nearest neighbor aggregates

## ... purpose

enhance each (stock, time) entry with features based on similar contexts using k-Nearest Neighbors (k-NN)

$$\text{log\_return}_t = \ln \left( \frac{\text{WAP}_t}{\text{WAP}_{t-1}} \right)$$

### > time-based nearest neighbors (time-NN)

- For each stock, construct vectors of recent *log\_return* values (e.g. last 10 buckets).
- Apply Euclidean distance to find k-nearest time windows within the same stock.
- Aggregate: mean\_ret\_3

### > stock-based nearest neighbors (stock-NN)

- For each time ID, normalize features like log\_wap across stocks
- Use k-NN to find similar stocks at the same time
- Aggregate: mean\_ret\_3

# ◎ LGBM

- > LightGBM is a high-performance gradient boosting framework based on decision trees
- > optimized for speed and memory usage, making it ideal for large tabular datasets
- > why it works?
  - captures nonlinear interactions
  - handles missing data efficiently
  - fast histogram-based learning
- > training objective
  - regression using mean squared error (MSE)

$$\hat{y} = \sum_{m=1}^M f_m(x), \quad f_m \in \text{Trees}$$

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

# ⌚ 1D-CNN

> 1D-CNN models sequences of time-series data (e.g. log-returns across previous time buckets) to detect local patterns and temporal structures

> why it works?

- efficient at capturing temporal locality
- learns spatial hierarchies with depth
- suited for sequence modeling

> structure

- input: shape (“n\_samples”, “seq\_len”, “n\_features”)
- layers: Conv1D → ReLU → MaxPooling → Flatten → Dense

$$h_t = \sigma \left( \sum_{k=1}^K w_k \cdot x_{t-k+1} + b \right)$$

# MLP

> Multi Layer Perceptron is a fully connected neural network that maps static tabular features to predictions using nonlinear transformations

> why it works?

- learns complex nonlinear relationships
- effective with engineered features
- simple architecture with high flexibility

> structure

- Input → Dense (ReLU) → Dense (ReLU) → Output (Linear)

For layers  $l = 1, 2, \dots, L$ :

$$a^{(l)} = \sigma(W^{(l)}a^{(l-1)} + b^{(l)})$$

Final output:

$$\hat{y} = W^{(L)}a^{(L-1)} + b^{(L)}$$



# Cross-validation techniques

... to evaluate model generalization and prevent overfitting by simulating real-time prediction scenarios

> time-ordered cross validation (time-series aware folds)

- unlike random shuffling, time-series data requires folds that respect temporal order
- for each fold:
  - - training is done on earlier time windows
  - - validation is performed on a later, unseen time segment

> advantages

- Mimics real-world forecasting, ensuring model learns only from the past.
- Prevents data leakage from future time buckets.
- Enhances robustness of model selection and hyperparameter tuning.

> Process

1. Sort data by `time_ord` (reconstructed chronological order).
2. Divide into  $K$  folds (e.g., 5-fold).
3. For fold  $k$ , train on time buckets " $< t_k$ "

$$\text{MSE}_k = \frac{1}{n_k} \sum_{i=1}^{n_k} (y_i - \hat{y}_i)^2$$

$$\text{Final MSE} = \frac{1}{K} \sum_{k=1}^K \text{MSE}_k$$

# ◎ ensemble and weights optimization

to improve prediction accuracy by combining the strengths of multiple diverse models—LightGBM, 1D-CNN, and MLP

## > ensemble strategy

- Model Outputs: each model predicts the target independently
- Stacking: combine outputs as weighted average

## > weight optimization

- grid search or optimization over validation set to minimize Mean Squared Error (MSE)
- constraints:
  - Non-negative weights
  - Sum to 1 for normalization

$$\hat{y}_{ensemble} = w_1 \cdot \hat{y}_{LGB} + w_2 \cdot \hat{y}_{CNN} + w_3 \cdot \hat{y}_{MLP}$$

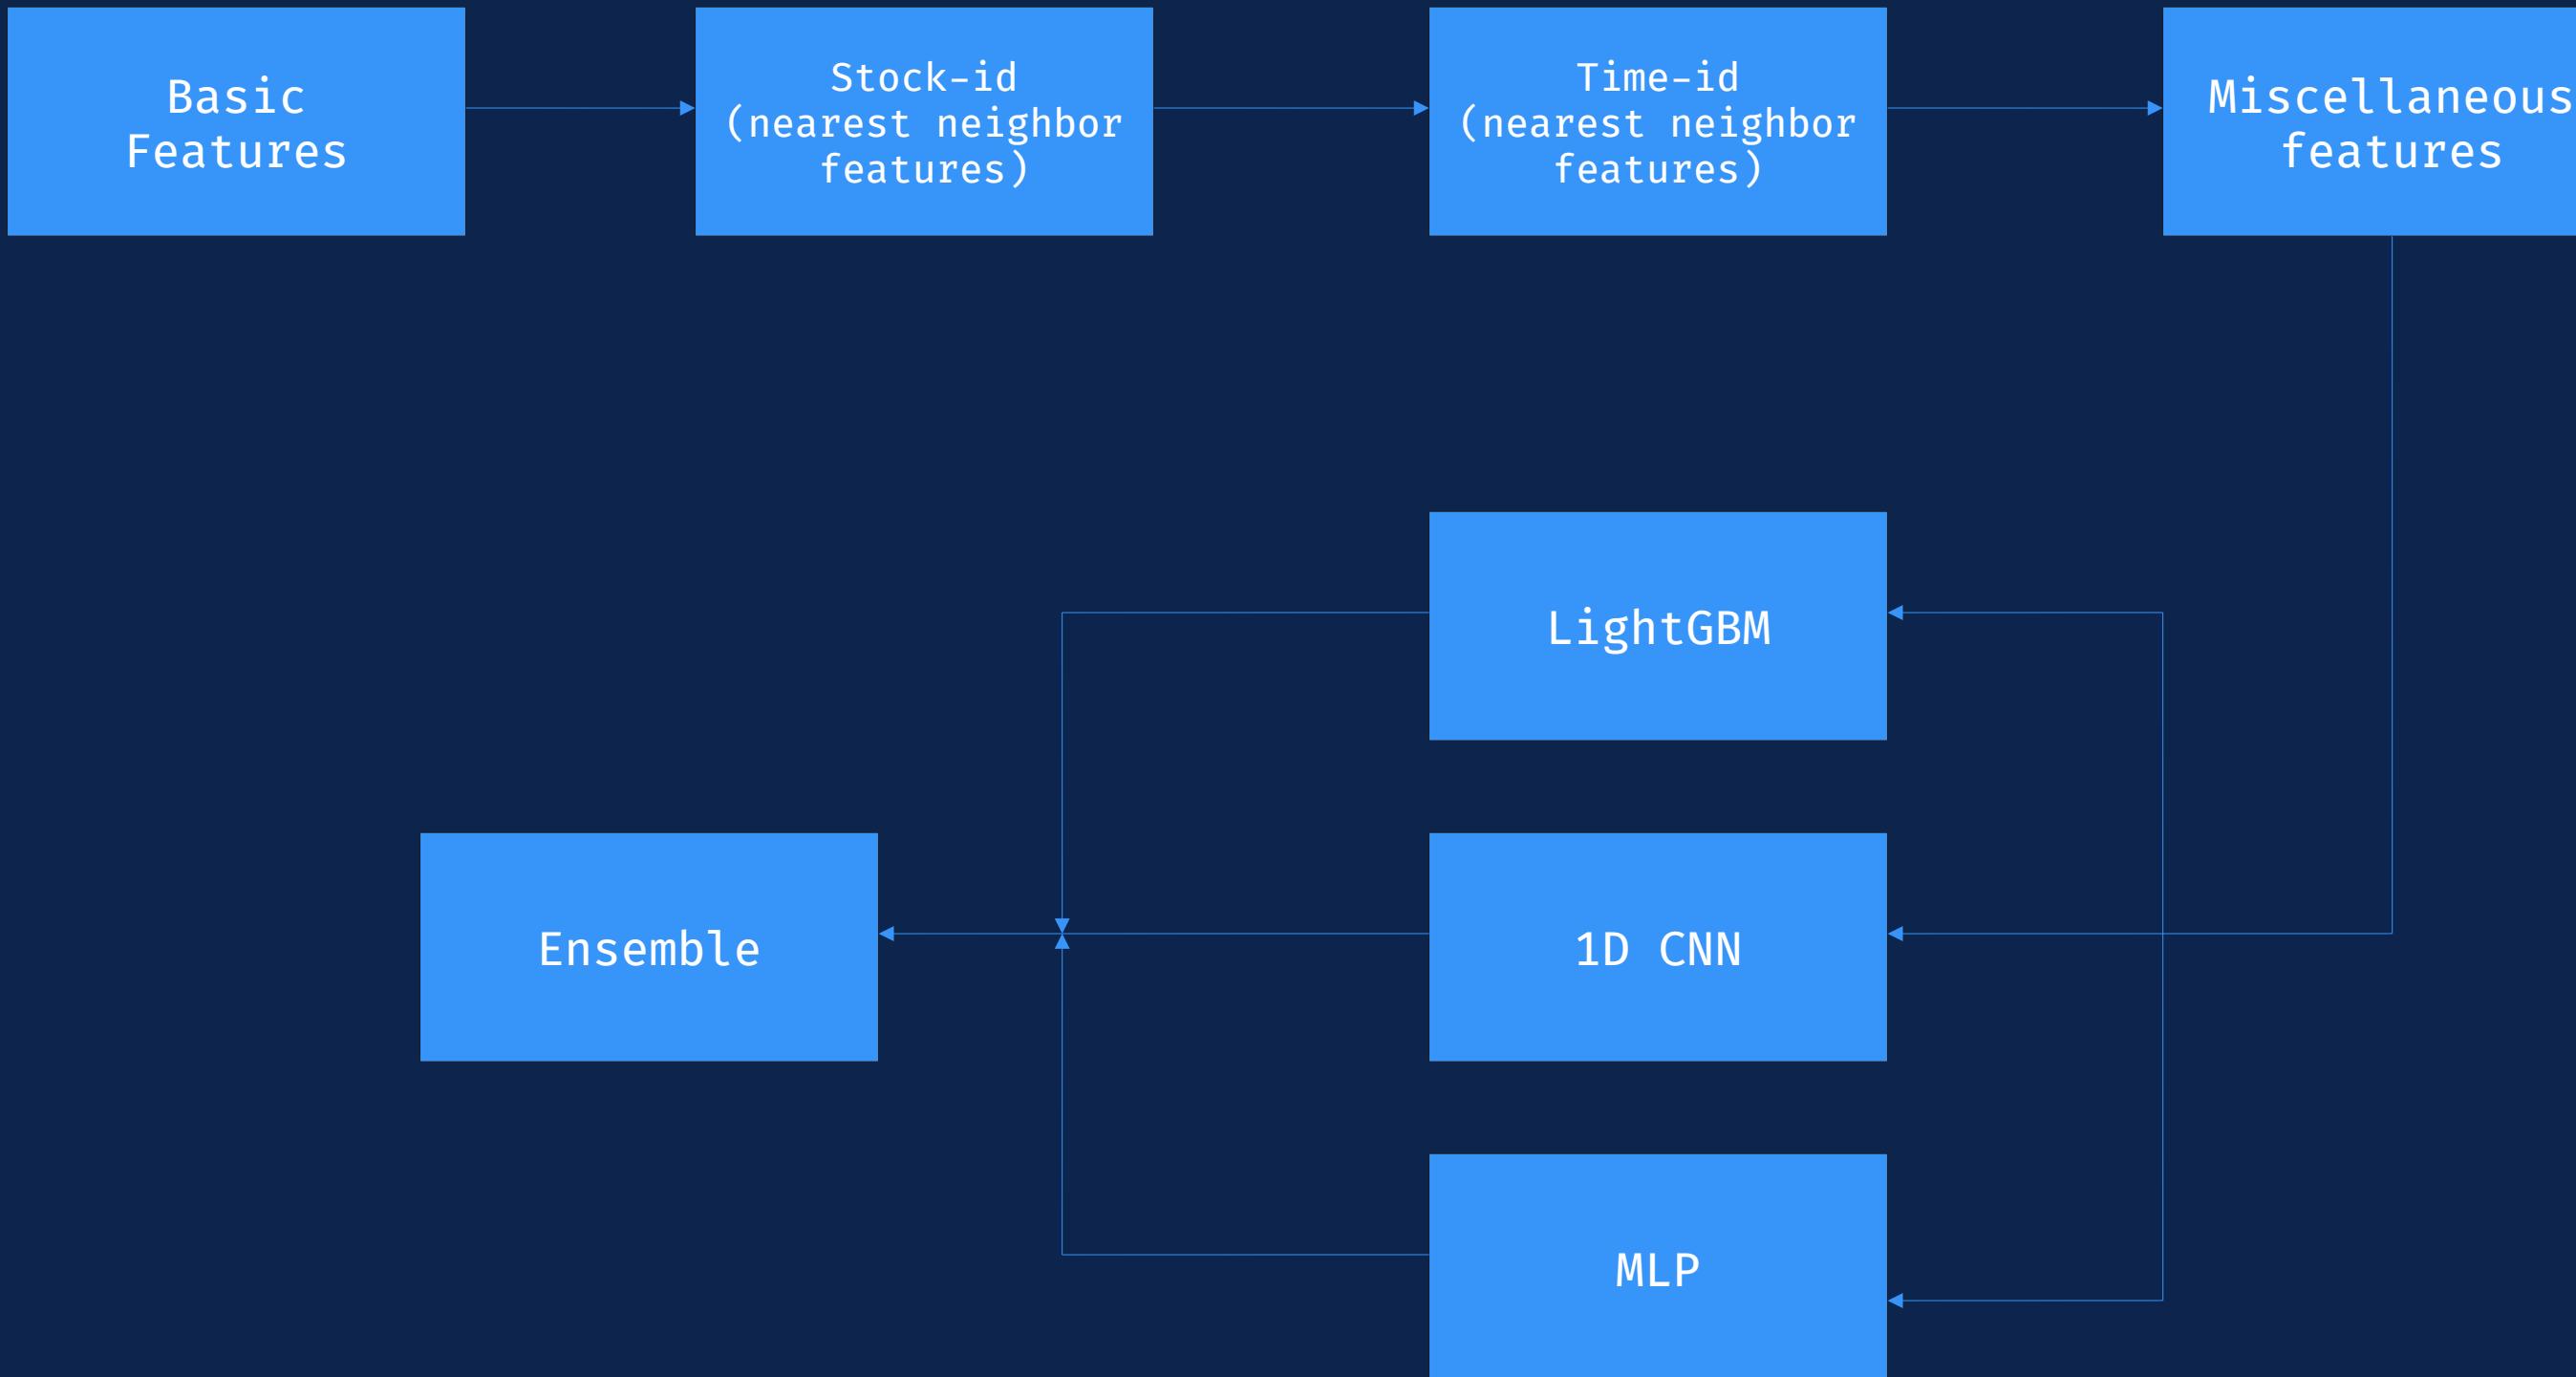
where  $w_1 + w_2 + w_3 = 1$

## > why ensemble works?

- Captures both linear/tabular (LGB), sequential (CNN), and dense relational (MLP) patterns
- Reduces individual model biases and variances

$$MSE(w) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_{ensemble}^{(i)})^2$$

# ⌚ workflow summary



# 🌀 results

Model	Fold 1 MSE	Fold 2 MSE	Fold 3 MSE	Fold 4 MSE	Fold 5 MSE	Avg MSE
LightGBM	.0.0172	0.0180	0.0169	0.0175	0.0178	0.0175
CNN	0.0165	0.0169	0.0167	0.0171	0.0168	0.0168
MLP	0.0181	0.0176	0.0179	0.0180	0.0177	0.0179
Ensemble	-	-	-	-	-	0.0162

thank  
y $\odot$ u

